

Student ID: 700752115

Student Name: Rumana Shaik

Neural Networks & Deep Learning – Assignment - 8

Github link: <https://github.com/ShaikRumana301/Neural-Network-DL-Assignment-8.git>

Video Link:

https://drive.google.com/file/d/1vyNwjGHVu_IZANuy3AK2WESaBa_fICRa/view?usp=sharing

- In class programming:

1. Add one more hidden layer to autoencoder

```
[1] from keras.layers import Input, Dense
    from keras.models import Model

    # this is the size of our encoded representations
    encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

    # this is our input placeholder
    input_img = Input(shape=(784,))
    # "encoded" is the encoded representation of the input
    encoded = Dense(encoding_dim, activation='relu')(input_img)
    # "decoded" is the lossy reconstruction of the input
    decoded = Dense(784, activation='sigmoid')(encoded)
    # this model maps an input to its reconstruction
    autoencoder = Model(input_img, decoded)
    # this model maps an input to its encoded representation
    autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy')
    from keras.datasets import mnist, fashion_mnist
    import numpy as np
    (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
    x_train = x_train.astype('float32') / 255.
    x_test = x_test.astype('float32') / 255.
    x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
    x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

    autoencoder.fit(x_train, x_train,
                    epochs=5,
                    batch_size=256,
                    shuffle=True,
                    validation_data=(x_test, x_test))
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>
29515/29515 [=====] - 0s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>
26421880/26421880 [=====] - 2s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>
5148/5148 [=====] - 0s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>
4422102/4422102 [=====] - 1s 0us/step
Epoch 1/5
235/235 [=====] - 4s 15ms/step - loss: 0.6972 - val_loss: 0.6970
Epoch 2/5
235/235 [=====] - 2s 10ms/step - loss: 0.6968 - val_loss: 0.6966
Epoch 3/5
235/235 [=====] - 4s 15ms/step - loss: 0.6964 - val_loss: 0.6962
Epoch 4/5
235/235 [=====] - 2s 10ms/step - loss: 0.6960 - val_loss: 0.6958
Epoch 5/5
235/235 [=====] - 2s 10ms/step - loss: 0.6957 - val_loss: 0.6955
<keras.src.callbacks.History at 0x7f3c3e93e230>

2. Do the prediction on the test data and then visualize one of the reconstructed version of that test data. Also, visualize the same test data before reconstruction using Matplotlib

✓ 18s

```
from keras.layers import Input, Dense
from keras.models import Model
from keras.datasets import mnist, fashion_mnist
import numpy as np
import matplotlib.pyplot as plt

encoding_dim = 32

input_img = Input(shape=(784,))

hidden_1 = Dense(256, activation='relu')(input_img)

encoded = Dense(encoding_dim, activation='relu')(hidden_1)

hidden_2 = Dense(256, activation='relu')(encoded)

# Define the output layer
decoded = Dense(784, activation='sigmoid')(hidden_2)

# Define the autoencoder model
autoencoder = Model(input_img, decoded)

# Compile the model
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics=['accuracy'])

# Load the fashion MNIST dataset
(x_train, _), (x_test, _) = fashion_mnist.load_data()
```

✓ 38s

```
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

history = autoencoder.fit(x_train, x_train,
                          epochs=5,
                          batch_size=256,
                          shuffle=True,
                          validation_data=(x_test, x_test))

decoded_imgs = autoencoder.predict(x_test)

# Visualize one of the reconstructed images
n = 10 # number of images to display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original test image
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)



    # Display reconstructed test image
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()

plt.plot(history.history['loss'])
```


```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper right')
plt.show()

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='lower right')
plt.show()
```

```
Epoch 1/5
235/235 [=====] - 8s 31ms/step - loss: 0.6934 - accuracy: 0.0040 - val_loss: 0.6933 - val_accuracy: 0.0037
Epoch 2/5
235/235 [=====] - 6s 25ms/step - loss: 0.6932 - accuracy: 0.0041 - val_loss: 0.6931 - val_accuracy: 0.0037
Epoch 3/5
235/235 [=====] - 6s 27ms/step - loss: 0.6931 - accuracy: 0.0040 - val_loss: 0.6930 - val_accuracy: 0.0037
Epoch 4/5
235/235 [=====] - 6s 26ms/step - loss: 0.6929 - accuracy: 0.0040 - val_loss: 0.6928 - val_accuracy: 0.0037
Epoch 5/5
235/235 [=====] - 8s 32ms/step - loss: 0.6927 - accuracy: 0.0040 - val_loss: 0.6926 - val_accuracy: 0.0036
313/313 [=====] - 1s 3ms/step
```



✓ [10]
38s



3. Repeat the question 2 on the denoising autoencoder

```
✓ [11] from keras.layers import Input, Dense
50s    from keras.models import Model

encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

input_img = Input(shape=(784,))
encoded = Dense(encoding_dim, activation='relu')(input_img)
# "decoded" is the lossy reconstruction of the input
decoded = Dense(784, activation='sigmoid')(encoded)
autoencoder = Model(input_img, decoded)
# this model maps an input to its encoded representation
autoencoder.compile(optimizer='adadelata', loss='binary_crossentropy')
from keras.datasets import fashion_mnist
import numpy as np
(x_train, _), (x_test, _) = fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

noise_factor = 0.5
x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

autoencoder.fit(x_train_noisy, x_train,
                epochs=10,
                batch_size=256,
                shuffle=True,
                validation_data=(x_test_noisy, x_test_noisy))
```

```
➡ Epoch 1/10
235/235 [=====] - 4s 14ms/step - loss: 0.6965 - val_loss: 0.6965
Epoch 2/10
235/235 [=====] - 2s 10ms/step - loss: 0.6963 - val_loss: 0.6962
Epoch 3/10
235/235 [=====] - 3s 11ms/step - loss: 0.6960 - val_loss: 0.6960
Epoch 4/10
235/235 [=====] - 2s 10ms/step - loss: 0.6958 - val_loss: 0.6958
Epoch 5/10
235/235 [=====] - 3s 12ms/step - loss: 0.6956 - val_loss: 0.6955
Epoch 6/10
235/235 [=====] - 3s 13ms/step - loss: 0.6953 - val_loss: 0.6953
Epoch 7/10
235/235 [=====] - 2s 10ms/step - loss: 0.6951 - val_loss: 0.6951
Epoch 8/10
235/235 [=====] - 2s 10ms/step - loss: 0.6949 - val_loss: 0.6949
Epoch 9/10
235/235 [=====] - 2s 11ms/step - loss: 0.6947 - val_loss: 0.6947
Epoch 10/10
235/235 [=====] - 3s 13ms/step - loss: 0.6945 - val_loss: 0.6945
<keras.src.callbacks.History at 0x7f3c2a33ec80>
```

4. plot loss and accuracy using the history object

```

from keras.layers import Input, Dense
from keras.models import Model
from keras.datasets import fashion_mnist
import numpy as np
import matplotlib.pyplot as plt

encoding_dim = 32

input_img = Input(shape=(784,))

encoded = Dense(encoding_dim, activation='relu')(input_img)

decoded = Dense(784, activation='sigmoid')(encoded)

autoencoder = Model(input_img, decoded)

# Compile the model
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics=['accuracy'])

# Load the fashion MNIST dataset
(x_train, _), (x_test, _) = fashion_mnist.load_data()

# Normalize the data and flatten the images
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

noise_factor = 0.5
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

noise_factor = 0.5
x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)

[12] x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

history = autoencoder.fit(x_train_noisy, x_train,
                          epochs=10,
                          batch_size=256,
                          shuffle=True,
                          validation_data=(x_test_noisy, x_test_noisy))

decoded_imgs = autoencoder.predict(x_test_noisy)

# Visualize one of the noisy test images
plt.figure(figsize=(20, 4))
n = 10
for i in range(n):
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test_noisy[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

# Visualize one of the reconstructed test images
for i in range(n):
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])


```

```
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper right')
plt.show()

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='lower right')
plt.show()
```

```
Epoch 1/10
235/235 [=====] - 4s 13ms/step - loss: 0.6980 - accuracy: 0.0014 - val_loss: 0.6977 - val_accuracy: 0.0012
Epoch 2/10
235/235 [=====] - 3s 13ms/step - loss: 0.6977 - accuracy: 0.0013 - val_loss: 0.6974 - val_accuracy: 0.0011
Epoch 3/10
235/235 [=====] - 3s 11ms/step - loss: 0.6974 - accuracy: 0.0013 - val_loss: 0.6971 - val_accuracy: 0.0011
Epoch 4/10
235/235 [=====] - 3s 11ms/step - loss: 0.6971 - accuracy: 0.0013 - val_loss: 0.6968 - val_accuracy: 0.0011
Epoch 5/10
235/235 [=====] - 3s 11ms/step - loss: 0.6968 - accuracy: 0.0013 - val_loss: 0.6966 - val_accuracy: 0.0012
Epoch 6/10
235/235 [=====] - 4s 15ms/step - loss: 0.6966 - accuracy: 0.0013 - val_loss: 0.6963 - val_accuracy: 0.0012
Epoch 7/10
```



Model Loss

